

03 Exercise Notebook 3

December 31, 2024

1 Exercise 3 - Brian Chen (bc604)

In this exercise, you will analyse a dataset obtained from the London transport system (TfL). The data is in a file called `tfl_readership.csv` (comma-separated-values format). As in Exercise 2, we will load and view the data using `pandas`.

```
[1]: # If you are running this on Google Colab, uncomment and run the following
      ↪ lines; otherwise ignore this cell
      # from google.colab import drive
      # drive.mount('/content/drive')
```

```
[2]: import math
      import numpy as np
      import matplotlib.pyplot as plt
      import pandas as pd
```

```
[3]: # Load data
      df_tfl = pd.read_csv('tfl_ridership.csv')
      # If running on Google Colab change path to '/content/drive/MyDrive/
      ↪ IB-Data-Science/Exercises/tfl_ridership.csv'

      df_tfl.head(13)
```

```
[3]:
```

	Year	Period	Start	End	Days	Bus cash (000s)	\
0	2000/01	P 01	01 Apr '00	29 Apr '00	29d	884	
1	2000/01	P 02	30 Apr '00	27 May '00	28d	949	
2	2000/01	P 03	28 May '00	24 Jun '00	28d	945	
3	2000/01	P 04	25 Jun '00	22 Jul '00	28d	981	
4	2000/01	P 05	23 Jul '00	19 Aug '00	28d	958	
5	2000/01	P 06	20 Aug '00	16 Sep '00	28d	984	
6	2000/01	P 07	17 Sep '00	14 Oct '00	28d	1001	
7	2000/01	P 08	15 Oct '00	11 Nov '00	28d	979	
8	2000/01	P 09	12 Nov '00	09 Dec '00	28d	971	
9	2000/01	P 10	10 Dec '00	06 Jan '01	28d	912	
10	2000/01	P 11	07 Jan '01	03 Feb '01	28d	943	
11	2000/01	P 12	04 Feb '01	03 Mar '01	28d	975	
12	2000/01	P 13	04 Mar '01	31 Mar '01	28d	974	

	Bus Oyster PAYG (000s)	Bus Contactless (000s) \
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0
5	0	0
6	0	0
7	0	0
8	0	0
9	0	0
10	0	0
11	0	0
12	0	0

	Bus One Day Bus Pass (000s)	Bus Day Travelcard (000s) ... \
0	210	231 ...
1	214	205 ...
2	209	221 ...
3	216	241 ...
4	225	248 ...
5	243	236 ...
6	205	216 ...
7	199	221 ...
8	184	212 ...
9	192	211 ...
10	193	186 ...
11	194	210 ...
12	186	204 ...

	Tube Contactless (000s)	Tube Day Travelcard (000s) \
0	0	655
1	0	605
2	0	650
3	0	708
4	0	730
5	0	702
6	0	639
7	0	668
8	0	640
9	0	631
10	0	556
11	0	617
12	0	584

	Tube Season Travelcard (000s)	Tube Other incl free (000s) \
0	1066	200

1	1168	217
2	1154	212
3	1196	214
4	1165	165
5	1164	151
6	1286	196
7	1298	220
8	1302	242
9	993	195
10	1259	234
11	1237	246
12	1262	266

	Tube Total (000s)	TfL Rail (000s)	Overground (000s)	DLR (000s)	\
0	2509	0	0	96	
1	2598	0	0	93	
2	2623	0	0	98	
3	2761	0	0	105	
4	2643	0	0	103	
5	2608	0	0	100	
6	2763	0	0	107	
7	2819	0	0	113	
8	2839	0	0	114	
9	2359	0	0	90	
10	2634	0	0	110	
11	2688	0	0	120	
12	2699	0	0	119	

	Tram (000s)	Air Line (000s)
0	45.8	0.0
1	46.5	0.0
2	47.1	0.0
3	50.8	0.0
4	50.3	0.0
5	49.2	0.0
6	48.8	0.0
7	51.5	0.0
8	54.0	0.0
9	55.3	0.0
10	50.1	0.0
11	50.5	0.0
12	47.7	0.0

[13 rows x 26 columns]

Each row of our data frame represents the average daily ridership over a 28/29 day period for various types of transport and tickets (bus, tube etc.). We have used the `.head()` command to display the top 13 rows of the data frame (corresponding to one year). Focusing on the “Tube

Total” column, notice the dip in ridership in row 9 (presumably due to Christmas/New Year’s), and also the slight dip during the summer (rows 4,5).

```
[4]: #df_tfl.sample(3) #random sample of 3 rows
df_tfl.tail(3) #last 3 rows
```

```
[4]:      Year Period      Start      End Days  Bus cash (000s)  \
242  2018/19    P 09  11 Nov '18  08 Dec '18  28d              0
243  2018/19    P 10  09 Dec '18  05 Jan '19  28d              0
244  2018/19    P 11  06 Jan '19  02 Feb '19  28d              0

      Bus Oyster PAYG (000s)  Bus Contactless (000s)  \
242                      1110                      1089
243                      1001                      949
244                      1036                      1075

      Bus One Day Bus Pass (000s)  Bus Day Travelcard (000s)  ...  \
242                      0                      41  ...
243                      0                      38  ...
244                      0                      30  ...

      Tube Contactless (000s)  Tube Day Travelcard (000s)  \
242                      1399                      249
243                      1110                      242
244                      1310                      204

      Tube Season Travelcard (000s)  Tube Other incl free (000s)  \
242                      1017                      334
243                      632                      259
244                      924                      305

      Tube Total (000s)  TfL Rail (000s)  Overground (000s)  DLR (000s)  \
242                      4221                      996                      557                      355
243                      3279                      750                      414                      270
244                      3809                      929                      517                      333

      Tram (000s)  Air Line (000s)
242          84.1          2.6
243          66.3          3.2
244          79.3          2.3
```

[3 rows x 26 columns]

The dataframe contains $N = 245$ counting periods (of 28/29 days each) from 1 April 2000 to 2 Feb 2019. We now define a numpy array consisting of the values in the 'Tube Total (000s)' column:

```
[5]: yvals = np.array(df_tfl['Tube Total (000s)'])
N = np.size(yvals)
```


```
xvals = np.linspace(1,N,N) #an array containing the values 1,2,...,N
```

We now have a time series consisting of points (x_i, y_i) , for $i = 1, \dots, N$, where y_i is the average daily tube rideship in counting period $x_i = i$.

1.1 3a) Plot the data in a scatterplot

```
[6]: #Your code for scatterplot here
```

```
plt.scatter(xvals, yvals, s=15)  
plt.show()
```



1.2 3b) Fit a linear model $f(x) = \beta_0 + \beta_1 x$ to the data

- Print the values of the regression coefficients β_0, β_1 determined using least-squares.
- Plot the fitted model and the scatterplot on the same plot.
- Compute and print the **MSE** and the R^2 coefficient for the fitted model.

All numerical outputs should be displayed to three decimal places.

```
[7]: #Your code here
```

```
X = np.column_stack((xvals**0, xvals))
```

```

beta = np.linalg.lstsq(X, yvals, rcond=None)[0]
fit = X.dot(beta)

print(f"beta_0 = {beta[0]:.3f}, beta_1 = {beta[1]:.3f}")

plt.plot(xvals, fit)
plt.scatter(xvals, yvals, s=10)
plt.show()


SSE = np.linalg.norm(yvals - fit) ** 2
MSE = SSE / N

fit_0 = np.array([np.mean(yvals) for _ in range(N)])
SSE_0 = np.linalg.norm(yvals - fit_0) ** 2
MSE_0 = SSE_0 / N

print(f"MSE = {MSE:.3f}, R^2 = {(1 - SSE/SSE_0):.3f}")

```

beta_0 = 2367.382, beta_1 = 5.939



MSE = 45323.636, R² = 0.796

1.3 3c) Plotting the residuals


- Plot the residuals on a scatterplot
- Also plot the residuals over a short duration and comment on whether you can discern any periodic components.


```
[8]: # Your code here
```

```
resids = yvals - fit

plt.scatter(xvals, resids, s=10)
plt.show()

plt.scatter(xvals[180:195], resids[180:195], s=10)
plt.show()
```





< Comment on periodic components here >

There doesn't appear to be any obvious periodic trends inside the residuals. Zooming out or zooming in has no significant difference in this regard.

1.3.1 3d) Periodogram

- Compute and plot the periodogram of the residuals. (Recall that the periodogram is the squared-magnitude of the DFT coefficients.)
- Identify the indices/frequencies for which the periodogram value exceeds **50%** of the maximum.

[9]: *# Your code to compute and plot the periodogram*


```
T = 1

pgram = np.abs(np.fft.fft(resids, N)/N)**2
indices = np.linspace(0, (N-1), num = N)
freqs_in_hz = np.fft.fftfreq(N)/T
freqs_in_rads = freqs_in_hz*2*math.pi

plt.subplot(121)
plt.plot(indices, pgram)
plt.xlabel('k')
```



```
plt.ylabel('Squared mag.  $|G_k|^2/N^2$ ')
plt.subplot(122)
plt.plot(freqs_in_hz, pgram)
plt.xlabel('frequency (1/mo)')
plt.show()
```



```
[10]: # Your code to identify the indices for which the periodogram value exceeds 50%
      ↪ of the maximum
```

```
top_indices = indices[(pgram > 0.5*np.max(pgram))]
print(top_indices)
```

```
[ 1.  38.  56.  57. 113. 132. 188. 189. 207. 244.]
```

1.4 3e) To the residuals, fit a model of the form

$$\beta_{1s} \sin(\omega_1 x) + \beta_{1c} \cos(\omega_1 x) + \beta_{2s} \sin(\omega_2 x) + \beta_{2c} \cos(\omega_2 x) + \dots + \beta_{Ks} \sin(\omega_K x) + \beta_{Kc} \cos(\omega_K x).$$

The frequencies $\omega_1, \dots, \omega_K$ in the model are those corresponding to the indices identified in Part 2c. (Hint: Each of the sines and cosines will correspond to one column in your X-matrix.)

- Print the values of the regression coefficients obtained using least-squares.

All numerical outputs should be displayed to three decimal places.

```
[11]: # Your code here

top_ws = 2 * np.pi * freqs_in_hz[(pgram > 0.5*np.max(pgram))]

XT = np.vstack([np.sin(w * xvals) for w in top_ws if w > 0] + [np.cos(w *
↪xvals) for w in top_ws if w > 0])
X = np.transpose(XT)

beta_sc = np.linalg.inv(XT.dot(X)).dot(XT).dot(resids)

print(f"beta_s = {beta_sc[0]:.3f}, beta_c = {beta_sc[1]:.3f}")
```

beta_s = -51.253, beta_c = 61.628

1.4.1 3f) The combined fit

- Plot the combined fit together with a scatterplot of the data
- Compute and print the final **MSE** and R^2 coefficient. Comment on the improvement over the linear fit.

The combined fit, which corresponds to the full model

$$f(x) = \beta_0 + \beta_1 x + \beta_{s1} \sin(\omega_1 x) + \beta_{c1} \cos(\omega_1 x) + \dots + \beta_{sk} \sin(\omega_k x) + \beta_{ck} \cos(\omega_k x),$$

can be obtained by adding the fits in parts 2b) and 2e).


```
[12]: # Your code here

fit_sc = X.dot(beta_sc)
fit_tot = fit + fit_sc

plt.plot(xvals, fit_tot)
plt.scatter(xvals, yvals, s=10)
plt.show()

SSE_tot = np.linalg.norm(yvals - fit_tot) ** 2
MSE_tot = SSE_tot / N

print(f"MSE = {MSE_tot:.3f}, R^2 = {(1 - SSE_tot/SSE_0):.3f}")
```



MSE = 20297.501, $R^2 = 0.908$

< Add comment on the improvement over the linear fit. >

Significant improvement of R^2 and significant reduction in MSE, showing that this fit is a lot better than the simple linear fit.